

МИНИСТЕРСТВО НАУКИ И ВЫСШЕГО ОБРАЗОВАНИЯ
РОССИЙСКОЙ ФЕДЕРАЦИИ
Федеральное государственное автономное образовательное учреждение
высшего образования
«КРЫМСКИЙ ФЕДЕРАЛЬНЫЙ УНИВЕРСИТЕТ им. В. И. ВЕРНАДСКОГО»
ФИЗИКО-ТЕХНИЧЕСКИЙ ИНСТИТУТ
Кафедра компьютерной инженерии и моделирования

ОТЧЕТ ПО ПРАКТИЧЕСКОМУ ЗАДАНИЮ №7
«Знакомство с Kubernetes»

Практическая работа
по дисциплине «Современные технологии программирования»
студента 1 курса группы ПИ-б-о-232
Халилов Асан Русланович
направления подготовки 09.03.04 «Программная инженерия»

Симферополь, 2024

Цель: Ознакомиться на практике с инструментом оркестрации контейнеризированных приложений Kubernetes.

Ход выполнения задания.

1,2,3,4,5

```
[node1 ~]$ kubeadm init --apiserver-advertise-address $(hostname -i) --pod-network-cidr 10.5.0.0/16
Initializing machine ID from random generator.
W0525 09:02:59.446824    520 initconfiguration.go:120] Usage of CRI endpoints without URL scheme is deprecated and can cause kubelet e
rrors in the future. Automatically prepending scheme "unix" to the "criSocket" with value "/run/docker/containerd/containerd.sock". Ple
ase update your configuration!
I0525 09:02:59.725916    520 version.go:256] remote version is much newer: v1.30.1; falling back to: stable-1.27
[init] Using Kubernetes version: v1.27.14
[preflight] Running pre-flight checks
[WARNING Swap]: swap is enabled; production deployments should disable swap unless testing the NodeSwap feature gate of the kub
elet
[preflight] The system verification failed. Printing the output from the verification:
KERNEL_VERSION: 4.4.0-210-generic
OS: Linux
CGROUPS_CPU: enabled
CGROUPS_CPUACCT: enabled
CGROUPS_CPUSET: enabled
CGROUPS_DEVICES: enabled
CGROUPS_FREEZER: enabled
CGROUPS_MEMORY: enabled
CGROUPS_PIDS: enabled
CGROUPS_HUGETLB: enabled
CGROUPS_BLKIO: enabled
[WARNING SystemVerification]: failed to parse kernel config: unable to load kernel module: "configs", output: "", err: exit sta
tus 1
[WARNING FileContent--proc-sys-net-bridge-bridge-nf-call-iptables]: /proc/sys/net/bridge/bridge-nf-call-iptables does not exist
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
W0525 09:03:00.166577    520 images.go:80] could not find officially supported version of etcd for Kubernetes v1.27.14, falling back t
o the nearest etcd version (3.5.7-0)
W0525 09:03:14.634321    520 checks.go:835] detected that the sandbox image "registry.k8s.io/pause:3.6" of the container runtime is in
consistent with that used by kubeadm. It is recommended that using "registry.k8s.io/pause:3.9" as the CRI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [kubernetes kubernetes.default kubernetes.default.svc kubernetes.default.svc.clu
ster.local node1] and IPs [10.96.0.1 192.168.0.18]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [localhost node1] and IPs [192.168.0.18 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key
[certs] Generating "sa" key and public key
[kubeconfig] Using kubeconfig folder "/etc/kubernetes"
[kubeconfig] Writing "admin.conf" kubeconfig file
[kubeconfig] Writing "kubelet.conf" kubeconfig file
[kubeconfig] Writing "controller-manager.conf" kubeconfig file
[kubeconfig] Writing "scheduler.conf" kubeconfig file
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Starting the kubelet
[control-plane] Using manifest folder "/etc/kubernetes/manifests"
[control-plane] Creating static Pod manifest for "kube-apiserver"
[control-plane] Creating static Pod manifest for "kube-controller-manager"
[control-plane] Creating static Pod manifest for "kube-scheduler"
[etcd] Creating static Pod manifest for local etcd in "/etc/kubernetes/manifests"
W0525 09:03:31.168847    520 images.go:80] could not find officially supported version of etcd for Kubernetes v1.27.14, falling back t
o the nearest etcd version (3.5.7-0)
[wait-control-plane] Waiting for the kubelet to boot up the control plane as static Pods from directory "/etc/kubernetes/manifests". Th
is can take up to 4m0s
[apiclient] All control plane components are healthy after 7.502469 seconds
[upload-config] Storing the configuration used in ConfigMap "kubeadm-config" in the "kube-system" Namespace
```

```
[kubelet] Creating a ConfigMap "kubelet-config" in namespace kube-system with the configuration for the kubelets in the cluster
[upload-certs] Skipping phase. Please see --upload-certs
[mark-control-plane] Marking the node node1 as control-plane by adding the labels: [node-role.kubernetes.io/control-plane node.kubernetes.io/exclude-from-external-load-balancers]
[mark-control-plane] Marking the node node1 as control-plane by adding the taints [node-role.kubernetes.io/control-plane:NoSchedule]
[bootstrap-token] Using token: 0qf246.7qk5vpgzyjt4ga4l
[bootstrap-token] Configuring bootstrap tokens, cluster-info ConfigMap, RBAC Roles
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to get nodes
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to post CSRs in order for nodes to get long term certificate credentials
[bootstrap-token] Configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] Configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy
```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
Alternatively, if you are the root user, you can run:
```

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 192.168.0.18:6443 --token 0qf246.7qk5vpgzyjt4ga4l \
--discovery-token-ca-cert-hash sha256:13814c75f4af5add9368d580dc7b8766703333c65f516c9e6f9a462e7a8da44
```

Waiting for api server to startup

Warning: resource daemonsets/kube-proxy is missing the kubect.kubernetes.io/last-applied-configuration annotation which is required by kubect apply. kubect apply should only be used on resources created declaratively by either kubect create --save-config or kubect apply. The missing annotation will be patched automatically.

daemonset.apps/kube-proxy configured

No resources found

```
[node1 ~]$ kubectl apply -f https://raw.githubusercontent.com/cloudnativelabs/kube-router/master/daemonset/kubeadm-kuberouter.yaml
configmap/kube-router-cfg created
daemonset.apps/kube-router created
serviceaccount/kube-router created
clusterrole.rbac.authorization.k8s.io/kube-router created
clusterrolebinding.rbac.authorization.k8s.io/kube-router created
```

6

```
[node1 ~]$ kubectl get nodes
NAME      STATUS    ROLES    AGE   VERSION
node1     Ready     control-plane   2m4s   v1.27.2
```

7

```
[node1 ~]$ kubectl get pods -A
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE
kube-system   coredns-5d78c9869d-jv5hx              1/1     Running   0           2m4s
kube-system   coredns-5d78c9869d-vzq5m              1/1     Running   0           2m4s
kube-system   etcd-node1                             1/1     Running   0           2m18s
kube-system   kube-apiserver-node1                   1/1     Running   0           2m10s
kube-system   kube-controller-manager-node1          1/1     Running   0           2m13s
kube-system   kube-proxy-pxcmg                       1/1     Running   0           2m4s
kube-system   kube-router-czxs                       1/1     Running   0           2m3s
kube-system   kube-scheduler-node1                   1/1     Running   0           2m16s
```

8

```
[node2 ~]$ kubeadm join 192.168.0.18:6443 --token 0qf246.7qk5vpgzyjt4ga4l \
> --discovery-token-ca-cert-hash sha256:13814c75f4af5add9368d580dc7b8766703333c65f516c9e6f9a462e7a8da44
Initializing machine ID from random generator.
W0525 09:06:38.813303 1254 initconfiguration.go:120] Usage of CRI endpoints without URL scheme is deprecated and can cause kubelet errors in the future. Automatically prepending scheme "unix" to the "criSocket" with value "/run/docker/containerd/containerd.sock". Please update your configuration!
[preflight] Running pre-flight checks
[WARNING Swap]: swap is enabled; production deployments should disable swap unless testing the NodeSwap feature gate of the kubelet
[preflight] The system verification failed. Printing the output from the verification:
KERNEL_VERSION: 4.4.0-210-generic
OS: Linux
CGROUPS_CPU: enabled
CGROUPS_CPUACCT: enabled
CGROUPS_CPUSET: enabled
CGROUPS_DEVICES: enabled
CGROUPS_FREEZER: enabled
CGROUPS_MEMORY: enabled
CGROUPS_PIDS: enabled
CGROUPS_HUGETLB: enabled
CGROUPS_BLKIO: enabled
```

```
[node3 ~]$ kubeadm join 192.168.0.18:6443 --token 0qf246.7qk5vpgzyjt4ga4l \
> --discovery-token-ca-cert-hash sha256:13814c75f4af5add9368d580dc7b8766703333c65f516c9e6f9a462e7a8da44
Initializing machine ID from random generator.
W0525 09:06:47.548387 1280 initconfiguration.go:120] Usage of CRI endpoints without URL scheme is deprecated and can cause kubelet e
rrors in the future. Automatically prepending scheme "unix" to the "criSocket" with value "/run/docker/containerd/containerd.sock". Ple
ase update your configuration!
[preflight] Running pre-flight checks
[WARNING Swap]: swap is enabled; production deployments should disable swap unless testing the NodeSwap feature gate of the kub
elet
[preflight] The system verification failed. Printing the output from the verification:
KERNEL_VERSION: 4.4.0-210-generic
OS: Linux
CGROUPS_CPU: enabled
CGROUPS_CPUACCT: enabled
CGROUPS_CPUSET: enabled
CGROUPS_DEVICES: enabled
CGROUPS_FREEZER: enabled
CGROUPS_MEMORY: enabled
CGROUPS_PIDS: enabled
CGROUPS_HUGETLB: enabled
CGROUPS_BLKIO: enabled
```

```
[node4 ~]$ kubeadm join 192.168.0.18:6443 --token 0qf246.7qk5vpgzyjt4ga4l \
> --discovery-token-ca-cert-hash sha256:13814c75f4af5add9368d580dc7b8766703333c65f516c9e6f9a462e7a8da44
Initializing machine ID from random generator.
W0525 09:06:55.606926 1299 initconfiguration.go:120] Usage of CRI endpoints without URL scheme is deprecated and can cause kubelet e
rrors in the future. Automatically prepending scheme "unix" to the "criSocket" with value "/run/docker/containerd/containerd.sock". Ple
ase update your configuration!
[preflight] Running pre-flight checks
[WARNING Swap]: swap is enabled; production deployments should disable swap unless testing the NodeSwap feature gate of the kub
elet
[preflight] The system verification failed. Printing the output from the verification:
KERNEL_VERSION: 4.4.0-210-generic
OS: Linux
CGROUPS_CPU: enabled
CGROUPS_CPUACCT: enabled
CGROUPS_CPUSET: enabled
CGROUPS_DEVICES: enabled
CGROUPS_FREEZER: enabled
CGROUPS_MEMORY: enabled
CGROUPS_PIDS: enabled
CGROUPS_HUGETLB: enabled
CGROUPS_BLKIO: enabled
```

```
[node5 ~]$ kubeadm join 192.168.0.18:6443 --token 0qf246.7qk5vpgzyjt4ga4l \
> --discovery-token-ca-cert-hash sha256:13814c75f4af5add9368d580dc7b8766703333c65f516c9e6f9a462e7a8da44
Initializing machine ID from random generator.
W0525 09:07:01.797187 1311 initconfiguration.go:120] Usage of CRI endpoints without URL scheme is deprecated and can cause kubelet e
rrors in the future. Automatically prepending scheme "unix" to the "criSocket" with value "/run/docker/containerd/containerd.sock". Ple
ase update your configuration!
[preflight] Running pre-flight checks
[WARNING Swap]: swap is enabled; production deployments should disable swap unless testing the NodeSwap feature gate of the kub
elet
[preflight] The system verification failed. Printing the output from the verification:
KERNEL_VERSION: 4.4.0-210-generic
OS: Linux
CGROUPS_CPU: enabled
CGROUPS_CPUACCT: enabled
CGROUPS_CPUSET: enabled
CGROUPS_DEVICES: enabled
CGROUPS_FREEZER: enabled
CGROUPS_MEMORY: enabled
CGROUPS_PIDS: enabled
CGROUPS_HUGETLB: enabled
CGROUPS_BLKIO: enabled
```

9

```
[node1 ~]$ kubectl get nodes
NAME      STATUS   ROLES    AGE     VERSION
node1     Ready    control-plane 3m41s   v1.27.2
node2     Ready    <none>    38s     v1.27.2
node3     Ready    <none>    29s     v1.27.2
node4     Ready    <none>    21s     v1.27.2
node5     Ready    <none>    15s     v1.27.2
```

```
[node1 ~]$ kubectl describe nodes node1
Name: node1
Roles: control-plane
Labels: beta.kubernetes.io/arch=amd64
        beta.kubernetes.io/os=linux
        kubernetes.io/arch=amd64
        kubernetes.io/hostname=node1
        kubernetes.io/os=linux
        node-role.kubernetes.io/control-plane=
        node.kubernetes.io/exclude-from-external-load-balancers=
Annotations: kubeadm.alpha.kubernetes.io/cri-socket: unix:///run/docker/containerd/containerd.sock
              node.alpha.kubernetes.io/ttl: 0
              volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Sat, 25 May 2024 09:03:36 +0000
Taints: node-role.kubernetes.io/control-plane:NoSchedule
Unschedulable: false
Lease:
  HolderIdentity: node1
  AcquireTime: <unset>
  RenewTime: Sat, 25 May 2024 09:07:31 +0000
Conditions:
  Type             Status  LastHeartbeatTime             LastTransitionTime             Reason                           Message
  ----             -
  MemoryPressure   False   Sat, 25 May 2024 09:04:07 +0000 Sat, 25 May 2024 09:03:34 +0000 KubeletHasSufficientMemory      kubelet has
  sufficient memory available
  DiskPressure     False   Sat, 25 May 2024 09:04:07 +0000 Sat, 25 May 2024 09:03:34 +0000 KubeletHasNoDiskPressure        kubelet has
  no disk pressure
  PIDPressure      False   Sat, 25 May 2024 09:04:07 +0000 Sat, 25 May 2024 09:03:34 +0000 KubeletHasSufficientPID         kubelet has
  sufficient PID available
  Ready            True    Sat, 25 May 2024 09:04:07 +0000 Sat, 25 May 2024 09:04:00 +0000 KubeletReady                    kubelet is
posting ready status
Addresses:
  InternalIP: 192.168.0.18
  Hostname: node1
Capacity:
  cpu: 8
  ephemeral-storage: 65504Mi
  hugepages-1Gi: 0
  hugepages-2Mi: 1364Mi
  memory: 32946972Ki
  pods: 110
Allocatable:
  cpu: 8
  ephemeral-storage: 65504Mi
  hugepages-1Gi: 0
  hugepages-2Mi: 1364Mi
  memory: 31550236Ki
  pods: 110
System Info:
  Machine ID: f22d41f7139549838106caa79eef572e
  System UUID: 0BE8ED01-977E-1847-BF47-D180FF4EEC73
  Boot ID: 7e47858a-6b84-421f-aa98-9609f7d294a0
  Kernel Version: 4.4.0-210-generic
  OS Image: CentOS Linux 7 (Core)
  Operating System: linux
  Architecture: amd64
  Container Runtime Version: containerd://1.6.21
  Kubelet Version: v1.27.2
  Kube-Proxy Version: v1.27.2
PodCIDR: 10.5.0.0/24
PodCIDRs: 10.5.0.0/24
Non-terminated Pods: (8 in total)
  Namespace             Name                               CPU Requests  CPU Limits  Memory Requests  Memory Limits  Age
  -----
  kube-system            coredns-5d78c9869d-jv5hx          100m (1%)    0 (0%)      70Mi (0%)       170Mi (0%)     3m45s
  kube-system            coredns-5d78c9869d-vzq5m          100m (1%)    0 (0%)      70Mi (0%)       170Mi (0%)     3m45s
  kube-system            etcd-node1                         100m (1%)    0 (0%)      100Mi (0%)       0 (0%)         3m59s
  kube-system            kube-apiserver-node1              250m (3%)    0 (0%)      0 (0%)           0 (0%)         3m51s
  kube-system            kube-controller-manager-node1     200m (2%)    0 (0%)      0 (0%)           0 (0%)         3m54s
  kube-system            kube-proxy-pxcmg                  0 (0%)       0 (0%)      0 (0%)           0 (0%)         3m45s
  kube-system            kube-router-czxvs                 250m (3%)    0 (0%)      250Mi (0%)       0 (0%)         3m44s
  kube-system            kube-scheduler-node1              100m (1%)    0 (0%)      0 (0%)           0 (0%)         3m57s
Allocated resources:
  (Total limits may be over 100 percent, i.e., overcommitted.)
  Resource           Requests    Limits
  -----
  cpu                 1100m (13%) 0 (0%)
  memory              490Mi (1%)  340Mi (1%)
  ephemeral-storage   0 (0%)      0 (0%)
  hugepages-1Gi       0 (0%)      0 (0%)
  hugepages-2Mi       0 (0%)      0 (0%)
Events:
  Type    Reason                     Age    From          Message
  ----    -
  Normal  Starting                   3m43s  kube-proxy    Starting kubelet.
  Normal  Starting                   4m6s   kubelet
  Warning  InvalidDiskCapacity        4m6s   kubelet       invalid capacity 0 on image filesystem
  Normal  NodeHasSufficientMemory    4m6s   (x8 over 4m6s) kubelet       Node node1 status is now: NodeHasSufficientMemory
  Normal  NodeHasNoDiskPressure      4m6s   (x8 over 4m6s) kubelet       Node node1 status is now: NodeHasNoDiskPressure
  Normal  NodeHasSufficientPID       4m6s   (x6 over 4m6s) kubelet       Node node1 status is now: NodeHasSufficientPID
  Normal  NodeAllocatableEnforced    4m6s   kubelet       Updated Node Allocatable limit across pods
  Normal  RegisteredNode              3m46s  node-controller Node node1 event: Registered Node node1 in Controller
```

11

```
[node1 ~]$ kubectl get pods -o wide -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
kube-system	coredns-5d78c9869d-jv5hx	1/1	Running	0	3m57s	10.5.0.4	node1	<none>		<none>	
kube-system	coredns-5d78c9869d-vzq5m	1/1	Running	0	3m57s	10.5.0.3	node1	<none>		<none>	
kube-system	etcd-node1	1/1	Running	0	4m11s	192.168.0.18	node1	<none>		<none>	
kube-system	kube-apiserver-node1	1/1	Running	0	4m3s	192.168.0.18	node1	<none>		<none>	
kube-system	kube-controller-manager-node1	1/1	Running	0	4m6s	192.168.0.18	node1	<none>		<none>	
kube-system	kube-proxy-66zqx	1/1	Running	0	47s	192.168.0.14	node5	<none>		<none>	
kube-system	kube-proxy-8jl7f	1/1	Running	0	61s	192.168.0.16	node3	<none>		<none>	
kube-system	kube-proxy-gqsvv	1/1	Running	0	53s	192.168.0.15	node4	<none>		<none>	
kube-system	kube-proxy-p74cl	1/1	Running	0	70s	192.168.0.17	node2	<none>		<none>	
kube-system	kube-proxy-pxcmg	1/1	Running	0	3m57s	192.168.0.18	node1	<none>		<none>	
kube-system	kube-router-8vb8k	1/1	Running	0	53s	192.168.0.15	node4	<none>		<none>	
kube-system	kube-router-czxvs	1/1	Running	0	3m56s	192.168.0.18	node1	<none>		<none>	
kube-system	kube-router-mtsl4	1/1	Running	0	61s	192.168.0.16	node3	<none>		<none>	
kube-system	kube-router-q868r	1/1	Running	0	47s	192.168.0.14	node5	<none>		<none>	
kube-system	kube-router-w7chf	1/1	Running	0	70s	192.168.0.17	node2	<none>		<none>	
kube-system	kube-scheduler-node1	1/1	Running	0	4m9s	192.168.0.18	node1	<none>		<none>	

12

```
[node1 ~]$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/website/master/content/en/examples/application/nginx-app.yaml
service/my-nginx-svc created
deployment.apps/my-nginx created
[node1 ~]$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
my-nginx-cbdccf466-2frn4	1/1	Running	0	27s	10.5.2.2	node3	<none>		<none>	
my-nginx-cbdccf466-jblcn	1/1	Running	0	27s	10.5.1.2	node2	<none>		<none>	
my-nginx-cbdccf466-pzqcm	1/1	Running	0	27s	10.5.4.2	node5	<none>		<none>	

```
[node1 ~]$
```

13

```
[node1 ~]$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	18m
my-nginx-svc	LoadBalancer	10.104.76.49	<pending>	80:31767/TCP	13m

```
[node1 ~]$
```

14

```
[node1 ~]$ curl 192.168.0.18:31767
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>
```

15

```
[node1 ~]$ kubectl drain node2 --ignore-daemonsets --delete-local-data
Flag --delete-local-data has been deprecated, This option is deprecated and will be deleted. Use --delete-emptydir-data.
node/node2 cordoned
Warning: ignoring DaemonSet-managed Pods: kube-system/kube-proxy-p74cl, kube-system/kube-router-w7chf
evicting pod default/my-nginx-cbdccf466-jblcn
pod/my-nginx-cbdccf466-jblcn evicted
node/node2 drained
[node1 ~]$ kubectl drain node3 --ignore-daemonsets --delete-local-data
Flag --delete-local-data has been deprecated, This option is deprecated and will be deleted. Use --delete-emptydir-data.
node/node3 cordoned
Warning: ignoring DaemonSet-managed Pods: kube-system/kube-proxy-8jl7f, kube-system/kube-router-mts14
evicting pod default/my-nginx-cbdccf466-2frn4
pod/my-nginx-cbdccf466-2frn4 evicted
node/node3 drained
[node1 ~]$ kubectl drain node5 --ignore-daemonsets --delete-local-data
Flag --delete-local-data has been deprecated, This option is deprecated and will be deleted. Use --delete-emptydir-data.
node/node5 cordoned
Warning: ignoring DaemonSet-managed Pods: kube-system/kube-proxy-66zqx, kube-system/kube-router-q868r
evicting pod default/my-nginx-cbdccf466-pzqcm
pod/my-nginx-cbdccf466-pzqcm evicted
node/node5 drained
[node1 ~]$ 
[node1 ~]$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
my-nginx-cbdccf466-4qks5	1/1	Running	0	72s	10.5.3.3	node4	<none>	<none>
my-nginx-cbdccf466-ftq6w	1/1	Running	0	62s	10.5.3.4	node4	<none>	<none>
my-nginx-cbdccf466-shld2	1/1	Running	0	81s	10.5.3.2	node4	<none>	<none>

```
[node1 ~]$
```

16,17

```
[node1 ~]$ kubectl delete -f https://raw.githubusercontent.com/kubernetes/website/master/content/en/examples/application/nginx-app.yaml
service "my-nginx-svc" deleted
deployment.apps "my-nginx" deleted
[node1 ~]$ kubectl get pods -o wide
No resources found in default namespace.
[node1 ~]$
```

Ответы на вопросы.

Как узнать IP-адрес ноды подключённой к кластеру по её имени? Вы можете использовать команду `kubectl get nodes -o wide` для получения информации о нодах, включая их IP-адреса.

Какой предварительный этап нужно выполнить, чтобы иметь возможность запустить приложение в кластере из исходников? Перед запуском приложения в кластере Kubernetes из исходного кода, вам нужно выполнить следующие шаги: Упаковать ваше приложение в контейнеры. Определить желаемое состояние вашего приложения с использованием манифестов. Загрузить ваш код приложения на платформу SCM, такую как GitHub. Использовать инструмент CI/CD для автоматизации развертывания вашего приложения.

Что нужно знать, чтобы подключить рабочую ноду к кластеру? Для подключения рабочей ноды к кластеру Kubernetes, вам нужно выполнить следующие шаги: Сгенерировать новый токен присоединения на мастер-ноде с помощью команды `kubeadm token create`. Присоединить новую рабочую ноду к кластеру с использованием токена присоединения и IP-адреса и порта мастер-ноды.

Может ли быть в одном кластере несколько мастер-нод? Да, в кластере Kubernetes может быть несколько мастер-нод. Это обеспечивает высокую доступность и защиту от отказов.

Может ли существовать кластер состоящий только из мастер ноды, которая одновременно является и рабочей нодой? Да, в теории, мастер-нода может также выполнять функции рабочей ноды и запускать рабочие нагрузки. Однако это не рекомендуется для производственных сред, главным образом из-за проблем производительности. В обучающих или ресурсно-ограниченных средах, вы можете иметь только одну ноду, которая выполняет функции и мастер-ноды, и рабочей ноды.

Вывод: Я ознакомился на практике с инструментом оркестрации контейнеризированных приложений Kubernetes.